

Salary Predictions Based on Job Descriptions

Part 1 - DEFINE

---- 1 Define the problem ----

The main problem of this project is to estimate salaries for new job posting based on historical job roles and their corresponding salaries

```
In [1]: #import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")

__author__ = "Bishnu Paudel"
__email__ = "ribishnu@gmail.com"
```

```
In [2]: pd.set_option('display.max_columns', 50)
```

Part 2 - DISCOVER

---- 2 Load the data ----

```
In [3]: #Load the data into a Pandas dataframe
train_salaries = pd.read_csv("/Users/bishnupaudel/Desktop/First_Portfolio_Python/tra
train_features = pd.read_csv("/Users/bishnupaudel/Desktop/First_Portfolio_Python/tra
test_features = pd.read_csv("/Users/bishnupaudel/Desktop/First_Portfolio_Python/test
```

---- 3 Clean the data ----

```
In [139]: #Look for duplicate data, invalid data (e.g. salaries <=0), or corrupt data and remo
```

```
In [4]: train_salaries.head()
```

```
Out[4]:
```

	jobId	salary
0	JOB1362684407687	130
1	JOB1362684407688	101
2	JOB1362684407689	137
3	JOB1362684407690	142
4	JOB1362684407691	163

```
In [5]: #checking size
print("The shape of the train_salaries is: ", train_salaries.shape)
```

```
print("The shape of the train_features is: ", train_features.shape)
print("The shape of the test_features is: ", test_features.shape)
```

```
The shape of the train_salaries is: (1000000, 2)
The shape of the train_features is: (1000000, 8)
The shape of the test_features is: (1000000, 8)
```

```
In [6]: test_features.columns, train_features.columns
```

```
Out[6]: (Index(['jobId', 'companyId', 'jobType', 'degree', 'major', 'industry',
               'yearsExperience', 'milesFromMetropolis'],
            dtype='object'),
        Index(['jobId', 'companyId', 'jobType', 'degree', 'major', 'industry',
               'yearsExperience', 'milesFromMetropolis'],
            dtype='object'))
```

```
In [7]: train_salaries.isnull().sum()
```

```
Out[7]: jobId      0
        salary      0
        dtype: int64
```

```
In [8]: train_salaries[train_salaries["salary"]==0].count()
```

```
Out[8]: jobId      5
        salary      5
        dtype: int64
```

```
In [9]: # there are five jobIDs whose salary are zero which is impossible. The data is signi
        ## it is better to replace those salary values with mode.
        train_salaries = train_salaries.replace(0, np.nan)
```

```
In [9]: train_salaries.min()
```

```
Out[9]: jobId      JOB1362684407687
        salary      0
        dtype: object
```

```
In [10]: train_features.isnull().sum()
```

```
Out[10]: jobId      0
        companyId    0
        jobType      0
        degree       0
        major        0
        industry     0
        yearsExperience 0
        milesFromMetropolis 0
        dtype: int64
```

```
In [11]: train_features.columns
```

```
Out[11]: Index(['jobId', 'companyId', 'jobType', 'degree', 'major', 'industry',
               'yearsExperience', 'milesFromMetropolis'],
            dtype='object')
```

```
In [12]: train_features.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000000 entries, 0 to 999999
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   jobId                 1000000 non-null object
1   companyId             1000000 non-null object
2   jobType               1000000 non-null object
3   degree               1000000 non-null object
```

```
4   major          1000000 non-null object
5   industry       1000000 non-null object
6   yearsExperience 1000000 non-null int64
7   milesFromMetropolis 1000000 non-null int64
dtypes: int64(2), object(6)
memory usage: 61.0+ MB
```

```
In [13]: test_features.isnull().sum()
```

```
Out[13]: jobId          0
companyId          0
jobType           0
degree            0
major             0
industry          0
yearsExperience    0
milesFromMetropolis 0
dtype: int64
```

```
In [14]: test_features.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000000 entries, 0 to 999999
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   jobId                 1000000 non-null object
1   companyId             1000000 non-null object
2   jobType               1000000 non-null object
3   degree                1000000 non-null object
4   major                 1000000 non-null object
5   industry              1000000 non-null object
6   yearsExperience        1000000 non-null int64
7   milesFromMetropolis    1000000 non-null int64
dtypes: int64(2), object(6)
memory usage: 61.0+ MB
```

```
In [15]: test_features.columns
```

```
Out[15]: Index(['jobId', 'companyId', 'jobType', 'degree', 'major', 'industry',
               'yearsExperience', 'milesFromMetropolis'],
              dtype='object')
```

```
In [16]: train_features.columns
```

```
Out[16]: Index(['jobId', 'companyId', 'jobType', 'degree', 'major', 'industry',
               'yearsExperience', 'milesFromMetropolis'],
              dtype='object')
```

```
In [17]: # Combining training datasets
#merge_df = pd.concat([train_salaries, train_features], axis=1)
#merge_df1 = merge_df.drop(["jobId", "companyId"], axis=1)

merge_df = pd.concat([train_features, test_features], axis=0)
merge_df = merge_df.drop(['jobId', 'companyId'], axis=1)
```

```
In [18]: merge_df.shape
```

```
Out[18]: (2000000, 6)
```

```
In [19]: merge_df.columns
```

```
Out[19]: Index(['jobType', 'degree', 'major', 'industry', 'yearsExperience',
               'milesFromMetropolis'],
              dtype='object')
```

```
In [20]: # Checking for duplicate values
merge_df.duplicated().sum()
```

Out[20]: 834305

```
In [21]: merge_df.isnull().sum()
```

```
Out[21]: jobType          0
degree          0
major          0
industry        0
yearsExperience  0
milesFromMetropolis  0
dtype: int64
```

```
In [170]: merge_df1 = merge_df1.fillna(merge_df["salary"].mean())
#merge_df = merge_df[(merge_df["salary"].notnull())]
#merge_df.isnull().sum()
```

```
In [171]: merge_df1.shape
```

Out[171]: (1000000, 7)

```
In [172]: merge_df1.isnull().sum()
```

```
Out[172]: salary          0
jobType          0
degree          0
major          0
industry        0
yearsExperience  0
milesFromMetropolis  0
dtype: int64
```

```
In [173]: ## Data looks healthy and no duplication of id in the data
```

---- 4 Explore the data (EDA) ----

```
In [22]: merge_df.describe()
```

```
Out[22]:
```

	yearsExperience	milesFromMetropolis
count	2.000000e+06	2.000000e+06
mean	1.199724e+01	4.952784e+01
std	7.212785e+00	2.888372e+01
min	0.000000e+00	0.000000e+00
25%	6.000000e+00	2.500000e+01
50%	1.200000e+01	5.000000e+01
75%	1.800000e+01	7.500000e+01
max	2.400000e+01	9.900000e+01

```
In [44]: # Exploring data types of merge data
merge_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2000000 entries, 0 to 999999
Data columns (total 6 columns):
```

```

#    Column                Dtype
---  -
0    jobType                object
1    degree                 object
2    major                  object
3    industry               object
4    yearsExperience        category
5    milesFromMetropolis    category
dtypes: category(2), object(4)
memory usage: 80.1+ MB

```

In [29]: *#JobId, companyId, jobType, degree, major, industry are category data and rest are n*

In [23]: `merge_df.columns`


Out[23]: Index(['jobType', 'degree', 'major', 'industry', 'yearsExperience',
'milesFromMetropolis'],
dtype='object')

In [24]: `merge_df.shape`

Out[24]: (2000000, 6)

In [33]: *# target variable*
`target = merge_df1["salary"]`
`plt.figure(figsize=(4,4))`
`plt.grid("white")`
`sns.distplot(target)`


Out[33]: <AxesSubplot:xlabel='salary', ylabel='Density'>



In [34]: *# target variable is following gaussian distribution and hence no outlier*

In [35]: `plt.figure(figsize=(8,4))`
`plt.grid("white")`
`sns.boxplot(target)`

Out[35]: <AxesSubplot:xlabel='salary'>



In []: *# Based on the boxplot, there are outliers which might need to consider to remove du*

In [177... `merge_df1['salary'].max(),merge_df1['salary'].min()`

Out[177... (301.0, 17.0)

```
In [178... mean = merge_df1['salary'].mean()
std = merge_df1['salary'].std()
x = mean-std*1.5
y = mean+std*1.5
merge_df_sal = merge_df1[(merge_df1['salary']>x) & (merge_df1['salary']<y)]
```


```
In [36]: corr = merge_df1[["salary", "yearsExperience", "milesFromMetropolis"]].corr()
corr
```

Out[36]:

	salary	yearsExperience	milesFromMetropolis
salary	1.000000	0.375012	-0.297686
yearsExperience	0.375012	1.000000	0.000673
milesFromMetropolis	-0.297686	0.000673	1.000000

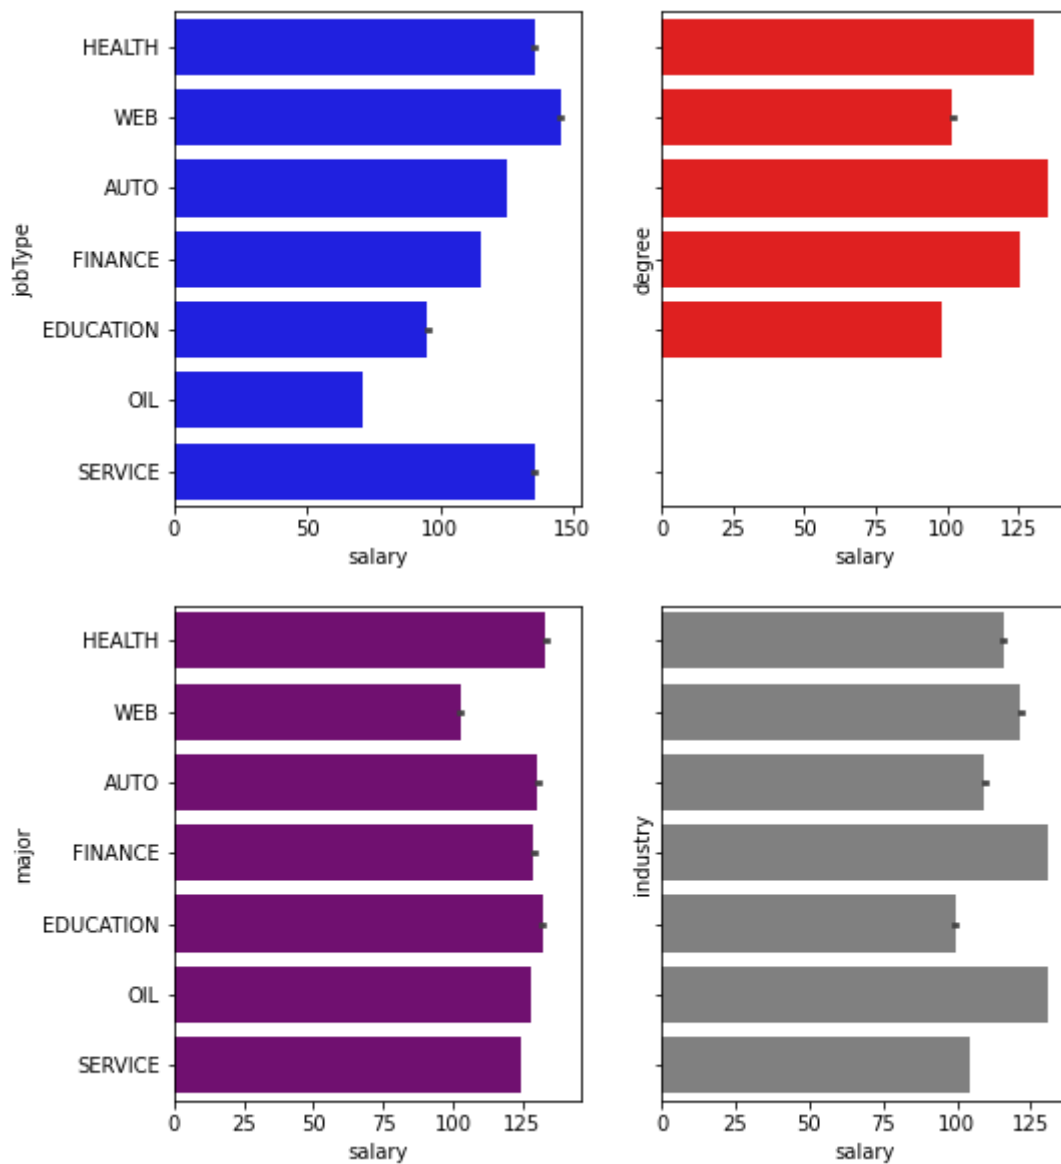
In [37]: `sns.heatmap(corr)`

Out[37]: <AxesSubplot:>



```
In [38]: #Bar plots showing how target variable depends upon category features
fig,axs = plt.subplots(2,2, figsize=(8,10), sharey=True)
sns.barplot(target, merge_df1["jobType"], color="blue", ax=axs[0,0])
sns.barplot(target, merge_df1["degree"],color="red", ax=axs[0,1])
sns.barplot(target, merge_df1["major"], color="purple", ax=axs[1,0])
sns.barplot(target, merge_df1["industry"],color="grey", ax=axs[1,1])
```

```
Out[38]: <AxesSubplot:xlabel='salary', ylabel='industry'>
```



---- 5 Establish a baseline ----

```
In [183... # Merge test data with train data
df = pd.concat([merge_df1, test_features_drop], axis=0)
df_sal = pd.concat([merge_df_sal, test_features_drop], axis=0)
```

```
In [184... df.shape
```

```
Out[184... (2000000, 7)
```

```
In [49]: merge_df.isnull().sum()
```

```
Out[49]: jobType          0
degree          0
major           0
industry        0
yearsExperience  0
milesFromMetropolis  0
dtype: int64
```

```
In [25]: merge_df["yearsExperience"].max()
```

```
Out[25]: 24
```

```
In [127... df_sal["yearsExperience"].max()
```


Out[127... 24

```
In [26]: merge_df["yearsExperience"] = pd.cut(merge_df["yearsExperience"], bins=[0,5,10,15,20],
include_lowest=True, labels=["low", "low-medium", "medium", "medium-high", "high"])
```

```
In [27]: merge_df["yearsExperience"]
```

```
Out[27]: 0      low-medium
1           low
2      low-medium
3      low-medium
4      low-medium
...
999995      medium
999996  medium-high
999997           low
999998      medium
999999  medium-high
Name: yearsExperience, Length: 2000000, dtype: category
Categories (5, object): ['low' < 'low-medium' < 'medium' < 'medium-high' < 'high']
```

```
In [28]: merge_df["milesFromMetropolis"].max()
```

Out[28]: 99

```
In [29]: merge_df["milesFromMetropolis"] = pd.cut(merge_df["milesFromMetropolis"], bins=[0,25],
labels=["very-near", "near", "medium", "far"])
```

```
In [31]: merge_df.head()
```

Out[31]:


	jobType	degree	major	industry	yearsExperience	milesFromMetropolis
0	CFO	MASTERS	MATH	HEALTH	low-medium	far
1	CEO	HIGH_SCHOOL	NONE	WEB	low	medium
2	VICE_PRESIDENT	DOCTORAL	PHYSICS	HEALTH	low-medium	near
3	MANAGER	DOCTORAL	CHEMISTRY	AUTO	low-medium	very-near
4	VICE_PRESIDENT	BACHELORS	PHYSICS	FINANCE	low-medium	very-near

```
In [30]: cat_features = ["yearsExperience", "milesFromMetropolis", "jobType", "degree", "major"]
dummies_df = pd.get_dummies(merge_df, columns=cat_features, drop_first=True)

dummies_df.tail()
```

Out[30]:

	yearsExperience_low-medium	yearsExperience_medium	yearsExperience_medium-high	yearsExperience_high
999995	0	1	0	0
999996	0	0	1	0
999997	0	0	0	0
999998	0	1	0	0
999999	0	0	1	0



```
In [ ]: cat_features = ["yearsExperience", "milesFromMetropolis", "jobType", "degree", "major"]
```

```
dummies_df_sal = pd.get_dummies(df_sal, columns=cat_features, drop_first=True)
```

```
In [31]: dummies_df.shape
```

```
Out[31]: (2000000, 32)
```

```
In [32]: dummies_df.isnull().sum()
```

```
Out[32]: yearsExperience_low-medium      0
yearsExperience_medium                0
yearsExperience_medium-high           0
yearsExperience_high                  0
milesFromMetropolis_near              0
milesFromMetropolis_medium            0
milesFromMetropolis_far               0
jobType_CFO                           0
jobType_CTO                           0
jobType_JANITOR                       0
jobType_JUNIOR                        0
jobType_MANAGER                       0
jobType_SENIOR                        0
jobType_VICE_PRESIDENT                0
degree_DOCTORAL                      0
degree_HIGH_SCHOOL                   0
degree_MASTERS                        0
degree_NONE                           0
major_BUSINESS                        0
major_CHEMISTRY                       0
major_COMPSCI                         0
major_ENGINEERING                     0
major_LITERATURE                      0
major_MATH                            0
major_NONE                            0
major_PHYSICS                         0
industry_EDUCATION                    0
industry_FINANCE                      0
industry_HEALTH                       0
industry_OIL                          0
industry_SERVICE                      0
industry_WEB                          0
dtype: int64
```

```
In [96]: #train_features1 = dummies_df.iloc[:868599,:]
#train_features1.shape
```

```
Out[96]: (868599, 33)
```

```
In [33]: train_features1 = dummies_df.iloc[:1000000,:]
train_features1.shape
```

```
Out[33]: (1000000, 32)
```

```
In [34]: test_features1 = dummies_df.iloc[1000000:, :]
test_features1.shape
```

```
Out[34]: (1000000, 32)
```

```
In [135... train_features2 = dummies_df_sal.iloc[:868599,:]
train_features1.shape
```

```
Out[135... (868599, 149)
```

```
In [128...
```

Out[128... (1000000, 33)

```
In [95]: merge_df_sal.shape
```

Out[95]: (868599, 7)

```
In [111... test_features2 = dummies_df.iloc[868599:,:]
test_features1.shape
```

Out[111... (1000000, 32)

```
In [39]: train_data = pd.concat([train_salaries, train_features1], axis=1)
train_data = train_data.drop('jobId', axis=1)
train_data.head()
```

Out[39]:

	salary	yearsExperience_low-medium	yearsExperience_medium	yearsExperience_medium-high	yearsExperience_high
0	130	1	0	0	
1	101	0	0	0	
2	137	1	0	0	
3	142	1	0	0	
4	163	1	0	0	

```
In [112... train_data.shape
```

Out[112... (1000000, 33)

```
In [98]: test_features2 = test_features1.drop('salary', axis=1)
test_features2.head()
```

Out[98]:

	yearsExperience_low-medium	yearsExperience_medium	yearsExperience_medium-high	yearsExperience_high	r
0	0	0	0	1	
1	0	0	1	0	
2	0	0	1	0	
3	0	1	0	0	
4	1	0	0	0	

```
In [47]: #df1 = df.sample(frac=0.5)
#df1.shape, df.shape
```

```
In [50]: train_data.isnull().sum()
```

Out[50]: salary 0
yearsExperience_low-medium 0
yearsExperience_medium 0
yearsExperience_medium-high 0
yearsExperience_high 0
milesFromMetropolis_near 0
milesFromMetropolis_medium 0
milesFromMetropolis_far 0

```

jobType_CFO                0
jobType_CTO                0
jobType_JANITOR            0
jobType_JUNIOR             0
jobType_MANAGER            0
jobType_SENIOR            0
jobType_VICE_PRESIDENT    0
degree_DOCTORAL           0
degree_HIGH_SCHOOL        0
degree_MASTERS             0
degree_NONE               0
major_BUSINESS            0
major_CHEMISTRY           0
major_COMPSCI             0
major_ENGINEERING         0
major_LITERATURE          0
major_MATH               0
major_NONE               0
major_PHYSICS             0
industry_EDUCATION        0
industry_FINANCE          0
industry_HEALTH           0
industry_OIL              0
industry_SERVICE          0
industry_WEB              0
dtype: int64

```

Segregation target and features for developmen of train and test datasets

```

In [72]: y = train_data["salary"]
X = train_data.drop('salary', axis=1)
X.shape, y.shape

```

```

Out[72]: ((1000000, 32), (1000000,))

```

Splitting data into train and test datasets with 80 to 20 ratio

```

In [73]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.2, random_state

```

```

In [52]: print("The shape of X_train is: ", X_train.shape)
print("The shape of X_test is: ", X_test.shape)
print("The shape of y_train is: ", y_train.shape)
print("The shape of y_test is: ", y_test.shape)

```

```

The shape of X_train is: (800000, 32)
The shape of X_test is: (200000, 32)
The shape of y_train is: (800000,)
The shape of y_test is: (200000,)

```

Building linear regression model as a base model

```

In [74]: from sklearn.linear_model import LinearRegression
model_lr = LinearRegression()
model_lr.fit(X_train, y_train)
y_pred_train = model_lr.predict(X_train)
y_pred_test = model_lr.predict(X_test)

```

```

In [75]: from sklearn.metrics import mean_squared_error
print("The MSE value of train is:", mean_squared_error(y_pred_train, y_train))

```

```
print("The MSE value of test is:", mean_squared_error(y_pred_test, y_test))
```

The MSE value of train is: 401.1362724752747
The MSE value of test is: 400.96545742447717

```
In [ ]: # The objective here is to reduce the MSE value, therefore feature engineering of d
# As the salary contains outliers I am going to remove those outliers shown by the b
```

Removing outliers from salary column

```
In [78]: mean = train_data['salary'].mean()
std = train_data['salary'].std()
x = mean-std*1.5
y = mean+std*1.5
merge_df_sal = train_data[(train_data['salary']>x) & (train_data['salary']<y)]
```

```
In [79]: merge_df_sal.shape
```

Out[79]: (868594, 33)

```
In [80]: train = merge_df_sal.iloc[:868594,:]
test = merge_df_sal.iloc[868594:,:]
train.shape, test.shape
```

Out[80]: ((868594, 33), (0, 33))

```
In [81]: test_y = train['salary']
train_X = train.drop('salary', axis=1)
train_X.shape, test_y.shape
```

Out[81]: ((868594, 32), (868594,))

Splitting data into train and test datasets

```
In [82]: from sklearn.model_selection import train_test_split
X_train1, X_test1, y_train1, y_test1 = train_test_split(train_X, test_y, random_stat
```

Building Linear Regression model after feature engineering done in salary

```
In [91]: from sklearn.linear_model import LinearRegression
model_lr_filt = LinearRegression()
model_lr_filt.fit(X_train1, y_train1)
y_pred_train_filt = model_lr.predict(X_train1)
y_pred_test_filt = model_lr.predict(X_test1)
```

```
In [92]: print("The MSE value of train is:", mean_squared_error(y_pred_train_filt, y_train1))
print("The MSE value of test is:", mean_squared_error(y_pred_test_filt, y_test1))
```

The MSE value of train is: 347.42401207552535
The MSE value of test is: 350.3922930349922

```
In [ ]: # After removing outliers the MSE value decrease significantly from 401 to 350. Now,
# I am going to try different machine learning models including RandomForest, Decisi
```

Building Random Forest Model

```
In [97]: from sklearn.ensemble import RandomForestRegressor
```

```

model_rf = RandomForestRegressor()
model_rf.fit(X_train1, y_train1)
y_train_rf_pred = model_rf.predict(X_train1)
y_test_rf_pred = model_rf.predict(X_test1)

```

```

In [102... print("The MSE value of train is:", mean_squared_error(y_train_rf_pred, y_train1))
            print("The MSE value of test dataset is:", mean_squared_error(y_test_rf_pred, y_test

```

The MSE value of train is: 280.8339257709302
The MSE value of test dataset is: 307.4691497808451

```

In [ ]: # As we can observe here that the MSE Value of test dataset is significantly decreases

```

Building XGBoost Model

```

In [104... #!pip install xgboost
import xgboost as xgb
model_xgb = xgb.XGBRegressor(n_estimators=100)
model_xgb.fit(X_train1, y_train1)
y_xgb_train_pred = model_xgb.predict(X_train1)
y_xgb_test_pred = model_xgb.predict(X_test1)

```

```

In [105... print("The MSE value of train is:", mean_squared_error(y_xgb_train_pred, y_train1))
            print("The MSE value of test dataset is:", mean_squared_error(y_xgb_test_pred, y_test

```

The MSE value of train is: 294.08317878976595
The MSE value of test dataset is: 297.6889212234916

```

In [ ]: # XGBoost resulted lowest MSE value for the testdata set

```

```

In [106... from sklearn.tree import DecisionTreeRegressor
model_tree = DecisionTreeRegressor()
model_tree.fit(X_train1, y_train1)
y_tree_pred_train = model_rf.predict(X_train1)
y_tree_pred_test = model_rf.predict(X_test1)

```

```

In [107... print("The MSE value of train is:", mean_squared_error(y_tree_pred_train, y_train1))
            print("The MSE value of test dataset is:", mean_squared_error(y_tree_pred_test, y_test

```

The MSE value of train is: 280.8339257709302
The MSE value of test dataset is: 307.4691497808451

```

In [ ]: # DecisionTree also yields similar MSE value in comparison to RF but slightly higher

```

```

In [ ]: # Overall, from above model fitting, XGBoost is the best model so far. I am going to
        # provide data set

```

Now, its time to predict the salary based on provide test_features

```

In [113... test_features1 = dummies_df.iloc[1000000:, :]
            test_features1.shape

```

Out[113... (1000000, 32)

```

In [121... y_xgb_test_pred = model_xgb.predict(test_features1)

```

```

In [125... #salary_pred = pd.DataFrame(y_xgb_test_pred)
            salary_pred = pd.DataFrame(y_xgb_test_pred)

```

```
salary_pred.head()
```

Out[125... 0

```
0 118.971931
1  94.759193
2 159.721466
3 107.175209
4 114.474030
```

In [117... test_features.head()

Out[117... jobld companyld jobType degree major industry yearsExperience m

```
0 JOB1362685407687 COMP33 MANAGER HIGH_SCHOOL NONE HEALTH 22
1 JOB1362685407688 COMP13 JUNIOR NONE NONE AUTO 20
2 JOB1362685407689 COMP10 CTO MASTERS BIOLOGY HEALTH 17
3 JOB1362685407690 COMP21 MANAGER HIGH_SCHOOL NONE OIL 14
4 JOB1362685407691 COMP36 JUNIOR DOCTORAL BIOLOGY OIL 10
```

In [126... test_features['xgb_salary_pred'] = salary_pred
test_features.head()

Out[126... jobld companyld jobType degree major industry yearsExperience m

```
0 JOB1362685407687 COMP33 MANAGER HIGH_SCHOOL NONE HEALTH 22
1 JOB1362685407688 COMP13 JUNIOR NONE NONE AUTO 20
2 JOB1362685407689 COMP10 CTO MASTERS BIOLOGY HEALTH 17
3 JOB1362685407690 COMP21 MANAGER HIGH_SCHOOL NONE OIL 14
4 JOB1362685407691 COMP36 JUNIOR DOCTORAL BIOLOGY OIL 10
```

In [130... salary_prediction = test_features[['jobId', 'xgb_salary_pred']]
salary_prediction

Out[130... jobld xgb_salary_pred

```
0 JOB1362685407687 118.971931
1 JOB1362685407688  94.759193
2 JOB1362685407689 159.721466
3 JOB1362685407690 107.175209
4 JOB1362685407691 114.474030
...
999995 JOB1362686407682 149.241287
999996 JOB1362686407683 106.678207
```

	jobld	xgb_salary_pred
999997	JOB1362686407684	65.381905
999998	JOB1362686407685	146.900024
999999	JOB1362686407686	114.610779

1000000 rows × 2 columns

In [134...

```
salary_prediction.to_csv("/Users/bishnupaudel/Desktop/First_Portfolio_Python/salary_
```